



# INTRODUCTION TO DEEP LEARNING

COURSE INSTRUCTOR: DR. M. UMAIR  
TOPIC: CONVOLUTIONAL NEURAL NETWORKS

## AGENDA

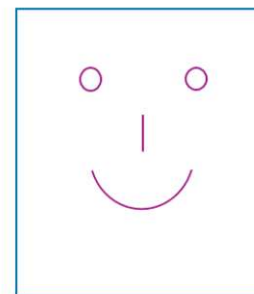
1. PRELIMINARIES
2. CONVOLUTIONAL NEURAL NETWORKS
3. CNN'S ARCHITECTURE
4. CODE EXAMPLE

## PRELIMINARIES

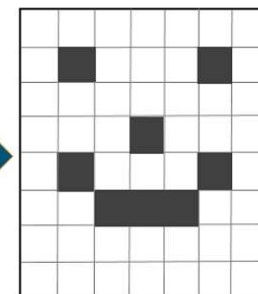
## PRELIMINARIES

### • Image Representation

- *Image* is represented in the form of an *array of pixel values*.



Real Image



Represented in the form of  
black and white pixels



|   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 1 | 0 | 0 | 0 | 1 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 1 | 0 | 0 | 0 |
| 0 | 1 | 0 | 0 | 0 | 1 | 0 |
| 0 | 0 | 1 | 1 | 1 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |

Image represented in the  
form of a matrix of numbers

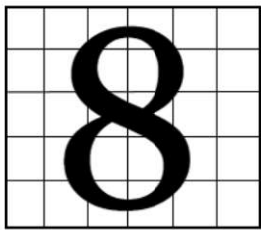
## PRELIMINARIES

### • Image Representation

- *Image* is represented in the form of an *array of pixel values*.



Real Image of the digit 8



Represented in the form of an array



|   |   |   |   |   |   |
|---|---|---|---|---|---|
| 0 | 0 | 1 | 1 | 0 | 0 |
| 0 | 1 | 0 | 0 | 1 | 0 |
| 0 | 0 | 1 | 1 | 0 | 0 |
| 0 | 1 | 0 | 0 | 1 | 0 |
| 0 | 0 | 1 | 1 | 0 | 0 |

Digit 8 represented in the form of pixels of 0's and 1's

## PRELIMINARIES

### • Image Representation – MNIST example



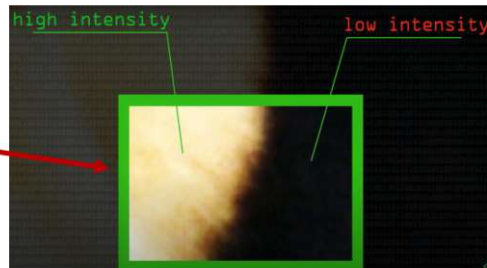
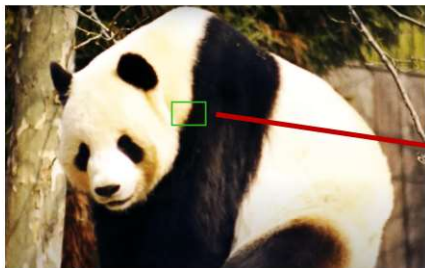
|    |     |     |     |     |     |     |     |     |     |     |     |     |     |     |    |
|----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|----|
| 0  | 2   | 15  | 0   | 0   | 11  | 10  | 0   | 0   | 0   | 0   | 9   | 9   | 0   | 0   | 0  |
| 0  | 0   | 0   | 4   | 60  | 158 | 236 | 255 | 255 | 177 | 95  | 61  | 32  | 0   | 0   | 29 |
| 0  | 10  | 16  | 113 | 238 | 255 | 244 | 245 | 243 | 250 | 249 | 255 | 222 | 10  | 0   | 0  |
| 0  | 14  | 170 | 255 | 255 | 244 | 254 | 255 | 253 | 245 | 255 | 249 | 253 | 251 | 13  | 1  |
| 2  | 68  | 255 | 228 | 255 | 251 | 254 | 211 | 141 | 116 | 122 | 215 | 251 | 238 | 255 | 49 |
| 13 | 217 | 243 | 255 | 155 | 33  | 226 | 62  | 2   | 0   | 10  | 13  | 232 | 255 | 255 | 36 |
| 16 | 229 | 252 | 254 | 48  | 12  | 0   | 0   | 7   | 7   | 0   | 70  | 237 | 252 | 235 | 62 |
| 6  | 141 | 245 | 255 | 217 | 25  | 11  | 9   | 3   | 0   | 115 | 236 | 243 | 255 | 137 | 0  |
| 0  | 87  | 252 | 250 | 248 | 215 | 60  | 0   | 1   | 171 | 252 | 255 | 248 | 174 | 6   | 0  |
| 0  | 13  | 113 | 255 | 255 | 245 | 255 | 182 | 181 | 248 | 252 | 242 | 208 | 36  | 0   | 19 |
| 1  | 0   | 5   | 117 | 251 | 255 | 241 | 255 | 247 | 255 | 241 | 162 | 17  | 0   | 7   | 0  |
| 0  | 0   | 0   | 4   | 58  | 251 | 255 | 246 | 254 | 253 | 255 | 120 | 11  | 0   | 1   | 0  |
| 0  | 0   | 4   | 97  | 255 | 255 | 255 | 248 | 252 | 255 | 244 | 255 | 182 | 10  | 0   | 4  |
| 0  | 22  | 206 | 252 | 246 | 251 | 241 | 100 | 24  | 113 | 255 | 245 | 255 | 194 | 9   | 0  |
| 0  | 111 | 255 | 242 | 255 | 158 | 24  | 0   | 0   | 6   | 39  | 255 | 232 | 230 | 56  | 0  |
| 0  | 218 | 251 | 250 | 137 | 7   | 11  | 0   | 0   | 2   | 62  | 255 | 250 | 125 | 3   | 0  |
| 0  | 173 | 255 | 255 | 101 | 9   | 20  | 0   | 13  | 3   | 13  | 182 | 251 | 245 | 61  | 0  |
| 0  | 107 | 251 | 241 | 255 | 230 | 98  | 55  | 19  | 118 | 217 | 248 | 253 | 255 | 52  | 4  |
| 0  | 18  | 146 | 250 | 255 | 247 | 255 | 255 | 245 | 249 | 255 | 240 | 255 | 129 | 0   | 5  |
| 0  | 0   | 23  | 113 | 215 | 255 | 250 | 248 | 255 | 255 | 248 | 248 | 118 | 14  | 12  | 0  |
| 0  | 0   | 6   | 1   | 0   | 52  | 153 | 233 | 255 | 252 | 147 | 37  | 0   | 0   | 4   | 1  |
| 0  | 0   | 5   | 5   | 0   | 0   | 0   | 0   | 0   | 14  | 1   | 0   | 6   | 6   | 0   | 0  |

|    |     |     |     |     |     |     |     |     |     |     |     |     |     |     |    |
|----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|----|
| 0  | 2   | 15  | 0   | 0   | 11  | 10  | 0   | 0   | 0   | 0   | 9   | 9   | 0   | 0   | 0  |
| 0  | 0   | 0   | 4   | 60  | 157 | 236 | 255 | 255 | 177 | 95  | 61  | 32  | 0   | 0   | 29 |
| 0  | 10  | 16  | 113 | 238 | 255 | 244 | 245 | 243 | 250 | 249 | 255 | 222 | 103 | 10  | 0  |
| 0  | 14  | 170 | 255 | 255 | 244 | 254 | 255 | 253 | 245 | 255 | 249 | 253 | 251 | 124 | 1  |
| 2  | 68  | 255 | 228 | 255 | 251 | 254 | 211 | 141 | 116 | 122 | 215 | 251 | 238 | 255 | 49 |
| 13 | 217 | 243 | 255 | 155 | 33  | 226 | 62  | 2   | 0   | 10  | 13  | 232 | 255 | 255 | 36 |
| 16 | 229 | 252 | 254 | 48  | 12  | 0   | 0   | 7   | 7   | 0   | 70  | 237 | 252 | 235 | 62 |
| 6  | 141 | 245 | 255 | 212 | 25  | 11  | 9   | 3   | 0   | 115 | 236 | 243 | 255 | 137 | 0  |
| 0  | 87  | 252 | 250 | 248 | 215 | 60  | 0   | 1   | 121 | 252 | 255 | 248 | 144 | 6   | 0  |
| 0  | 13  | 113 | 255 | 255 | 245 | 255 | 182 | 181 | 248 | 252 | 242 | 208 | 36  | 0   | 19 |
| 1  | 0   | 5   | 117 | 251 | 255 | 241 | 255 | 247 | 255 | 241 | 162 | 17  | 0   | 7   | 0  |
| 0  | 0   | 0   | 4   | 58  | 251 | 255 | 246 | 254 | 253 | 255 | 120 | 11  | 0   | 1   | 0  |
| 0  | 0   | 4   | 97  | 255 | 255 | 255 | 248 | 252 | 255 | 244 | 255 | 182 | 10  | 0   | 4  |
| 0  | 22  | 206 | 252 | 246 | 251 | 241 | 100 | 24  | 113 | 255 | 245 | 255 | 194 | 9   | 0  |
| 0  | 111 | 255 | 242 | 255 | 158 | 24  | 0   | 0   | 6   | 39  | 255 | 232 | 230 | 56  | 0  |
| 0  | 218 | 251 | 250 | 137 | 7   | 11  | 0   | 0   | 2   | 62  | 255 | 250 | 125 | 3   | 0  |
| 0  | 173 | 255 | 255 | 101 | 9   | 20  | 0   | 13  | 3   | 13  | 182 | 251 | 245 | 61  | 0  |
| 0  | 107 | 251 | 241 | 255 | 230 | 98  | 55  | 19  | 118 | 217 | 248 | 253 | 255 | 52  | 4  |
| 0  | 18  | 146 | 250 | 255 | 247 | 255 | 255 | 245 | 249 | 255 | 240 | 255 | 129 | 0   | 5  |
| 0  | 0   | 23  | 113 | 215 | 255 | 250 | 248 | 255 | 255 | 248 | 248 | 118 | 14  | 12  | 0  |
| 0  | 0   | 6   | 1   | 0   | 52  | 153 | 233 | 255 | 252 | 147 | 37  | 0   | 0   | 4   | 1  |
| 0  | 0   | 5   | 5   | 0   | 0   | 0   | 0   | 0   | 14  | 1   | 0   | 6   | 6   | 0   | 0  |

## PRELIMINARIES

### • Edge Detection

- *Sharp change* in intensity or color.
- *High value* indicates *steep change*.
- *Low value* indicates *shallow change*.



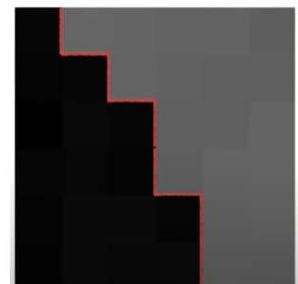
## PRELIMINARIES

### • Edge Detection

- It's a techniques which *identifies significant transitions* in *pixel intensity or color*, indicating the boundaries of objects within an image.

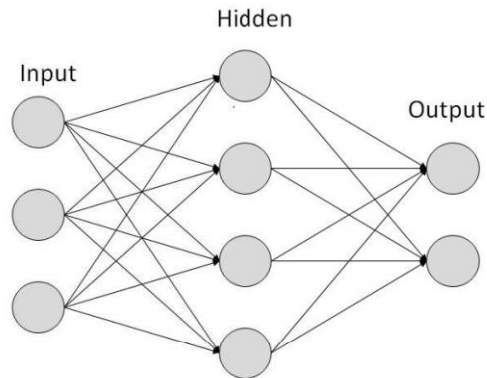
|    |    |    |    |    |    |
|----|----|----|----|----|----|
| 12 | 90 | 89 | 86 | 87 | 82 |
| 10 | 12 | 88 | 85 | 83 | 84 |
| 9  | 15 | 12 | 84 | 84 | 88 |
| 12 | 14 | 10 | 82 | 88 | 89 |
| 11 | 17 | 16 | 12 | 88 | 90 |
| 10 | 16 | 15 | 17 | 89 | 88 |

|    |    |    |    |    |    |
|----|----|----|----|----|----|
| 12 | 90 | 89 | 86 | 87 | 82 |
| 10 | 12 | 88 | 85 | 83 | 84 |
| 9  | 15 | 12 | 84 | 84 | 88 |
| 12 | 14 | 10 | 82 | 88 | 89 |
| 11 | 17 | 16 | 12 | 88 | 90 |
| 10 | 16 | 15 | 17 | 89 | 88 |



## PRELIMINARIES

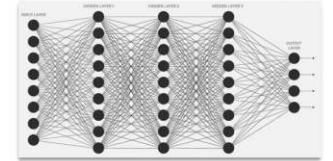
### • Basic Neural Network



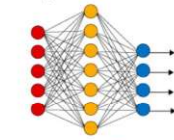
## PRELIMINARIES

### • Deep Neural Network

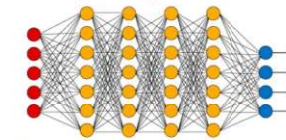
- *Deep* neural *network* consists of *more hidden layers*.
- Each *input* is *connected* to the *hidden layer* and the neural *network* will *decide* the *connections*.



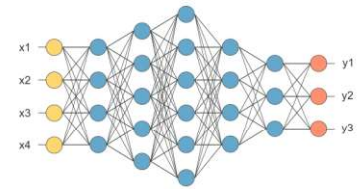
Simple Neural Network



Deep Learning Neural Network



● Input Layer ● Hidden Layer ● Output Layer

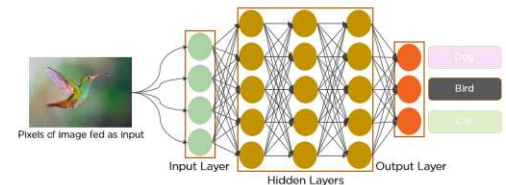


## CONVOLUTIONAL NEURAL NETWORKS

## CONVOLUTIONAL NEURAL NETWORKS

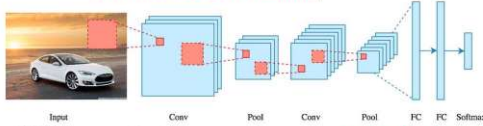
### • Convolutional Neural Network

- A **CONVOLUTIONAL NEURAL NETWORK** (ConvNet or CNN) is a *type of deep neural network architecture*.
- It is widely used for
  - Image classification
  - Object detection
  - Semantic segmentation
  - Natural language processing etc.

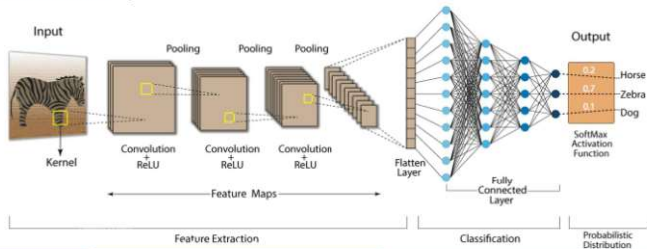


# CONVOLUTIONAL NEURAL NETWORKS

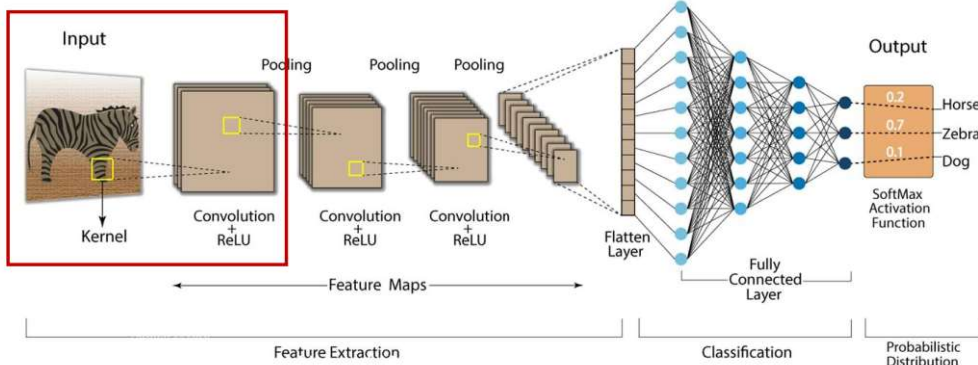
- All *CNN* models follow a *similar architecture*, as shown in the figure below.



- More specifically, CNNs typically consist of *convolutional*, *pooling*, *flatten*, and *fully connected* layers.



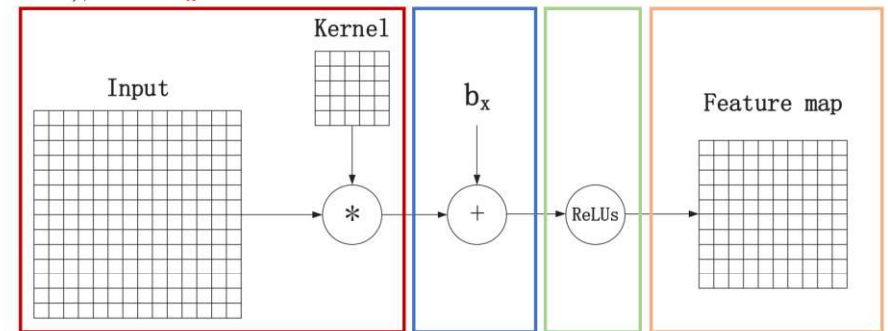
## CNN ARCHITECTURE CONVOLUTION LAYER



## CONVOLUTION LAYER

### • Overview

- Following is the *overall process* of *convolution layer* (in the simplified case), where  $b_x$  is the *bias*.





## CONVOLUTION LAYER

### • Convolution

- The *main idea* behind CNNs is to use *convolutional* layers.
- Convolution is a specialized type of *linear operation* used for *feature extraction*, where a small array of numbers, called a *kernel*, is applied across the *input*, which is an array of numbers, called a *tensor*. The results is a *feature map*.

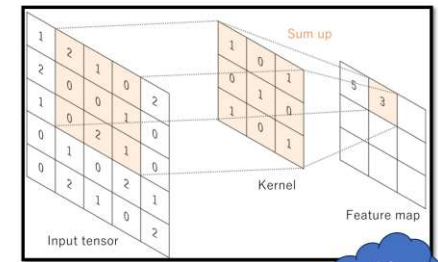
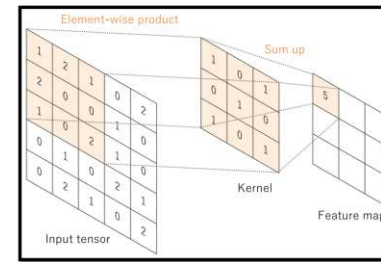
$$y_j = \sum k_{ij} * x_i$$

where  $*$  denotes the convolution operation,  $y_j$  is the  $j$ th feature map of output,  $k_{ij}$  is the trainable convolutional kernel (also called filter), and  $x_i$  is the  $i$ th input.

## CONVOLUTION LAYER

### • Convolution - Example

$$y_j = \sum k_{ij} * x_i$$



reduce the number of parameters

## CONVOLUTION LAYER

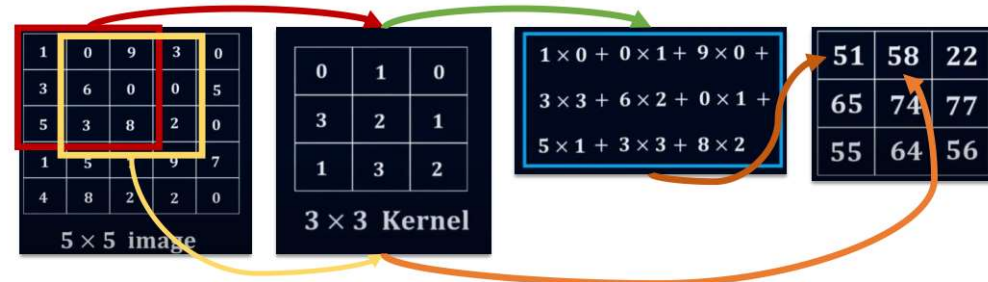
### • Kernels (filters) & Feature Maps

- The *convolutional layers* apply a *set of kernels (filters)* to the input data, which *results* in a *set of feature maps* that capture the features of the input data.
- The *number* of the *feature map* is the *same* as the *depth* of the convolution *kernel*.
- The *filters* are *learned during the training process*, and they are *used* to *detect patterns and features* in the input data.

## CONVOLUTION LAYER

### • Kernels or Filters

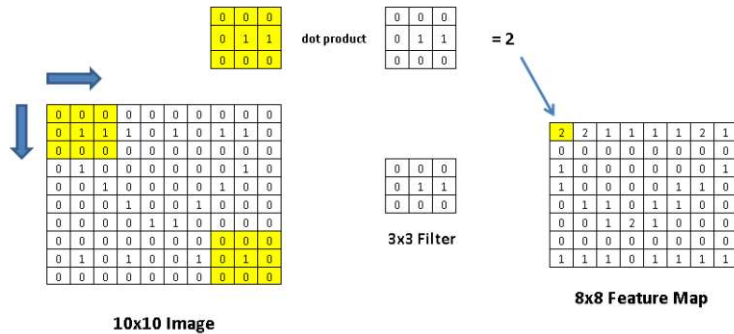
- Kernels (filters)* *reduces* the *dimension* and *extracts* the *features*.



## CONVOLUTION LAYER

### • Kernels or Filters

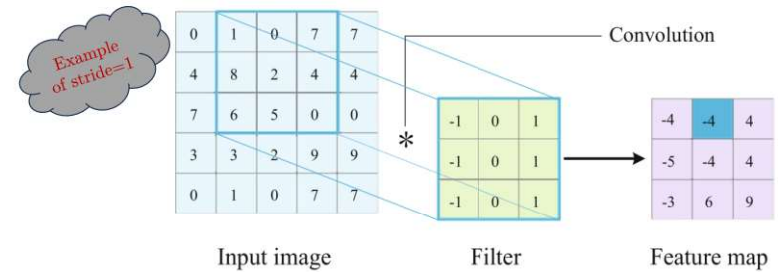
- *Kernels (filters) reduces the dimension and extracts the features.*



## CONVOLUTION LAYER

### • Strides

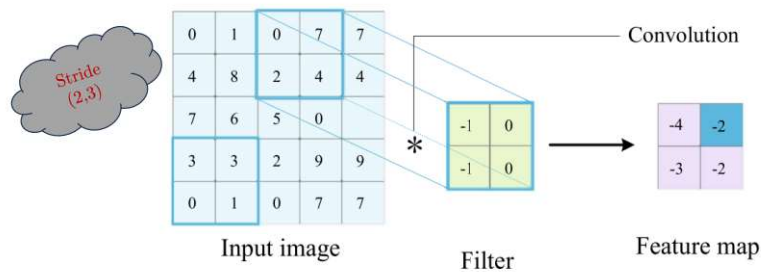
- **STRIDE** is the *number of rows and columns* that the convolution *kernel slides over* the *input matrix* in order from left to right and top to bottom, starting from the top left of the input matrix.



## CONVOLUTION LAYER

### • Strides

- Another *example* is a *stride of 2 in the horizontal direction and 3 in the vertical direction.*

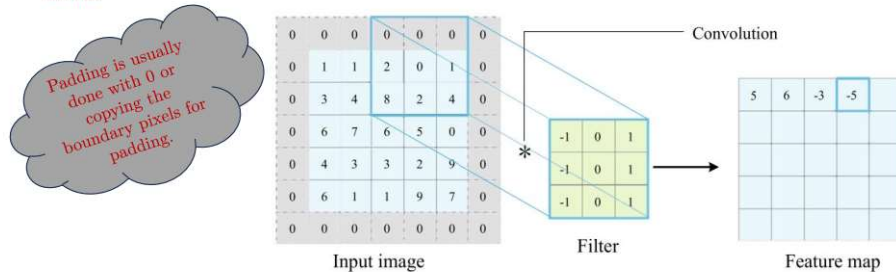


*The output data size, computational complexity, and feature extraction capability can all be impacted by the stride.*

## CONVOLUTION LAYER

### • Padding

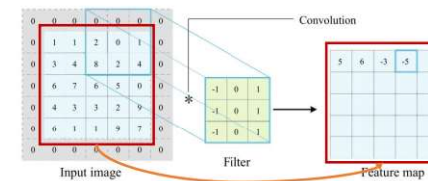
- PADDING** is the process of *adding* a certain number of *pixels* to the *edges* of the *input data* so that the *size* of the *output data* can *match* the *input data*.



## CONVOLUTION LAYER

### • Padding

- Padding* makes it easier for the convolution kernel to *learn* the *information surrounding* the *input image*.
- For instance, when the  $5 \times 5 \times 1$  image is *reinforced* into a  $7 \times 7 \times 1$  and applied to the  $3 \times 3 \times 1$  kernel over it, *the computed matrix* is shown to be of dimensions  $5 \times 5 \times 1$ .



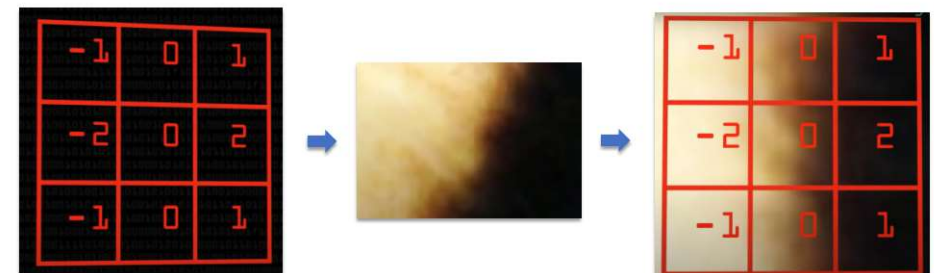
*If the same convolution procedure is done without padding, the output might have a smaller-sized image.*

*Consequently, a  $5 \times 5 \times 1$  image will be converted to a  $3 \times 3 \times 1$  image.*

## CONVOLUTION LAYER

### • Extracting Features

- Edge detection



*Negative value indicates light to dark.*

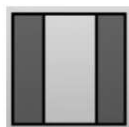
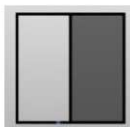
## CONVOLUTION LAYER

### • Extracting Features

- Extracting vertical edge

$$\begin{bmatrix} 10 & 10 & 10 & 0 & 0 & 0 \\ 10 & 10 & 10 & 0 & 0 & 0 \\ 10 & 10 & 10 & 0 & 0 & 0 \\ 10 & 10 & 10 & 0 & 0 & 0 \\ 10 & 10 & 10 & 0 & 0 & 0 \\ 10 & 10 & 10 & 0 & 0 & 0 \end{bmatrix} * \begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix} = \begin{bmatrix} 0 & 30 & 30 & 0 \\ 0 & 30 & 30 & 0 \\ 0 & 30 & 30 & 0 \\ 0 & 30 & 30 & 0 \\ 0 & 30 & 30 & 0 \\ 0 & 30 & 30 & 0 \end{bmatrix}$$

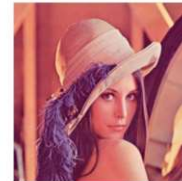
Positive value indicates light to dark.



## CONVOLUTION LAYER

### • Extracting Features

- Horizontal and vertical edges



Horizontal Edges



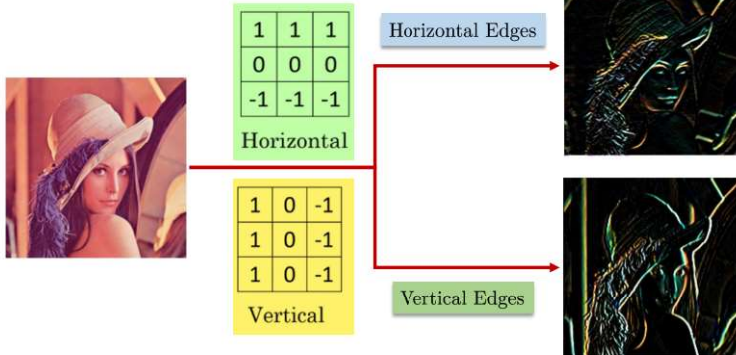
Vertical Edges



## CONVOLUTION LAYER

### • Extracting Features

- Horizontal and vertical edges



## CONVOLUTION LAYER

### • Extracting Features

- Some commonly used kernels for edge detection.

$$\mathbf{G}_x = \begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix} * \mathbf{A}$$

$$\mathbf{G}_y = \begin{bmatrix} +1 & +2 & +1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} * \mathbf{A}$$

SOBEL FILTER

$$\mathbf{G}_x = \begin{bmatrix} -3 & 0 & 3 \\ -10 & 0 & 10 \\ -3 & 0 & 3 \end{bmatrix} * \mathbf{I}$$

$$\mathbf{G}_y = \begin{bmatrix} -3 & -10 & -3 \\ 0 & 0 & 0 \\ 3 & 10 & 3 \end{bmatrix} * \mathbf{I}$$

SCHAAR FILTER



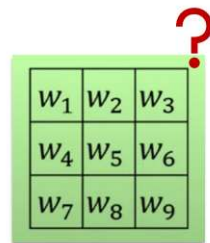
## CONVOLUTION LAYER

### • Kernel values, CNN perspective

- Which kernel values to choose?

|   |   |   |   |   |   |
|---|---|---|---|---|---|
| 3 | 0 | 1 | 2 | 7 | 4 |
| 1 | 5 | 8 | 9 | 3 | 1 |
| 2 | 7 | 2 | 5 | 1 | 3 |
| 0 | 1 | 3 | 1 | 7 | 8 |
| 4 | 2 | 1 | 6 | 2 | 8 |
| 2 | 4 | 5 | 2 | 3 | 9 |

A complex image

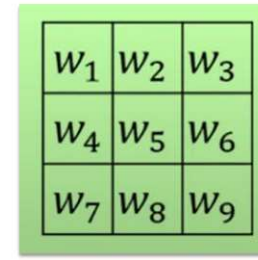


Kernel

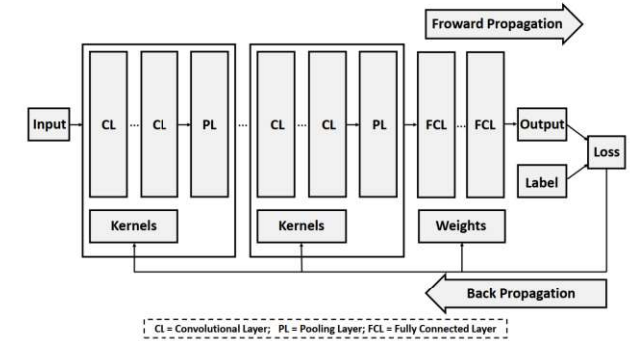
## CONVOLUTION LAYER

### • Kernel values, CNN perspective

- Which kernel values to choose?

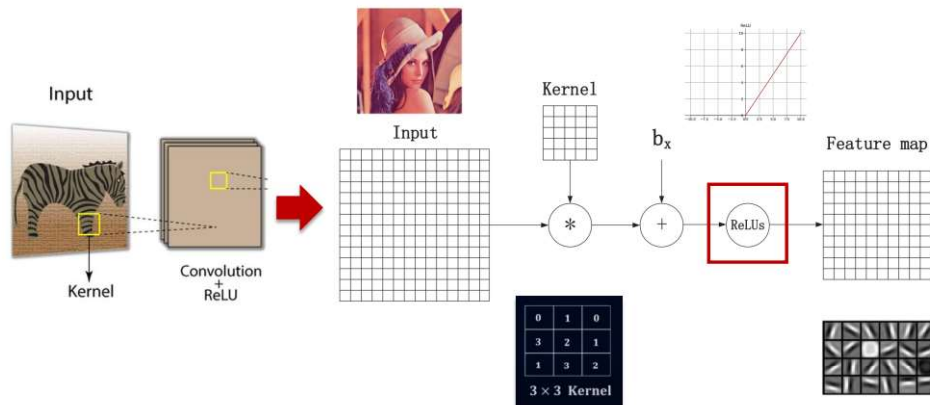


Learning these values makes CNN more robust as compared to computer vision



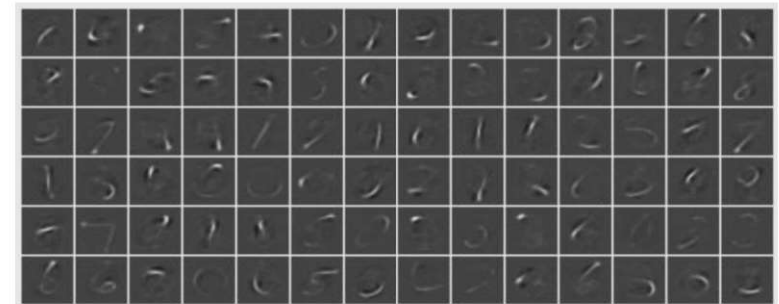
## CONVOLUTION LAYER

### • ReLU Activation



## CONVOLUTION LAYER

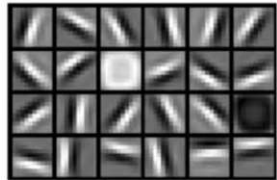
### • First CNN layer output



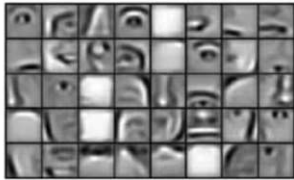
Activations taken from the *first convolutional layer* of a simplistic deep CNN, after training on the *MNIST dataset* of handwritten digits. If you look carefully, you can see that the *network has successfully picked up on characteristics* unique to specific numeric digits.

## CONVOLUTION LAYER

- Layer by layer output of a CNN



EARLIER LAYERS  
Detect edges

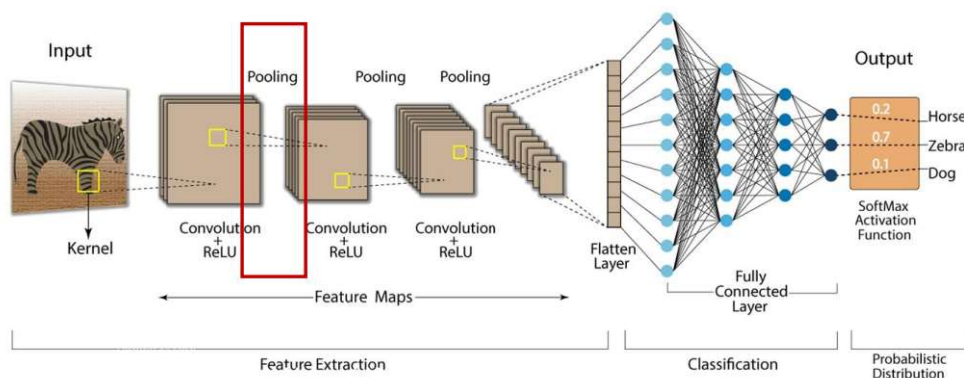


LATER LAYERS  
Detects part of object



EVEN LATER LAYERS  
Detects complete object

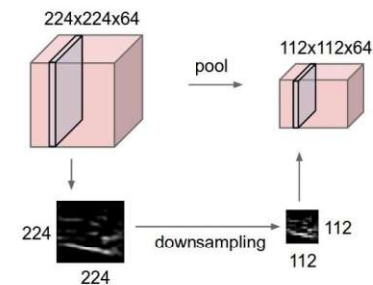
## CNN ARCHITECTURE POOLING LAYER



## POOLING LAYER

- Introduction

- The **POOLING LAYERS** are used to *down-sample* the feature maps, which helps to *reduce the number of parameters* in the model and *prevent overfitting*.



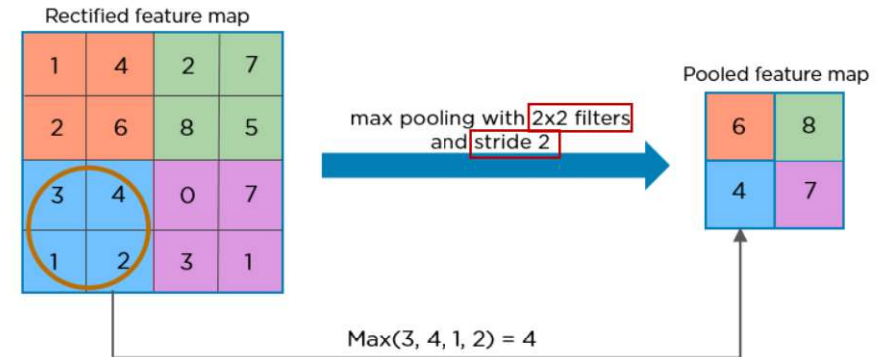
## POOLING LAYER

### • Introduction

- The *pooling procedure*, like the convolution process, can be thought of as a pooling *function without weights*, in which the *input feature* mapping group is divided into many *regions* and each area is pooled to *yield a value as a generalization* of this region.
- There are *several types of pooling layers*, including *max* pooling, *min* pooling, *average* pooling, and *sum* pooling etc.

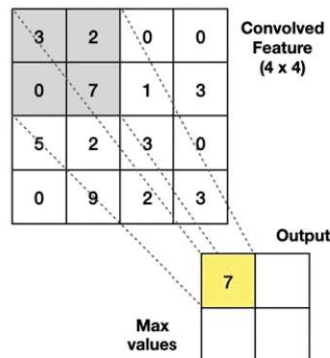
## POOLING LAYER

### • How it Works?



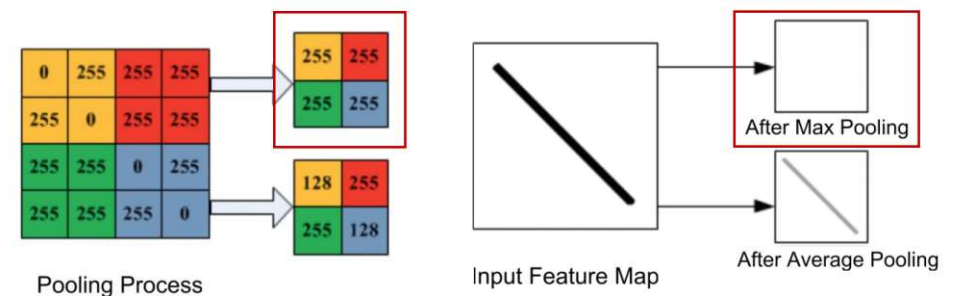
## POOLING LAYER

### • Max Pooling - Example



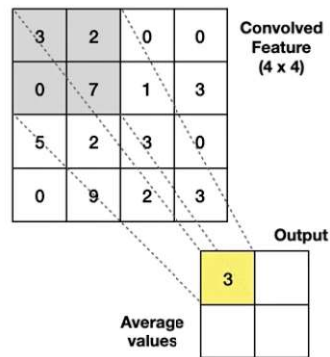
## POOLING LAYER

### • Max Pooling Result



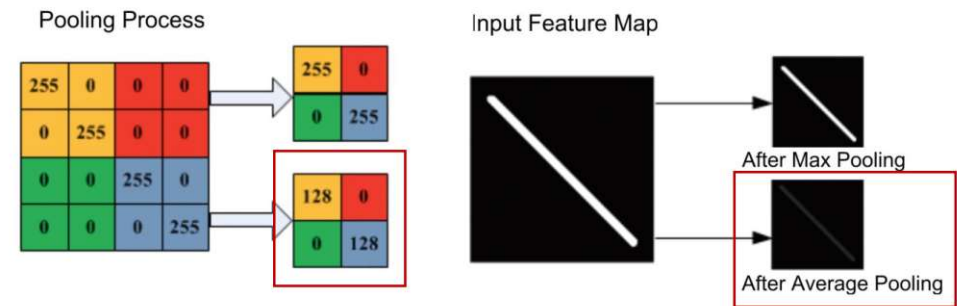
## POOLING LAYER

### • Average Pooling - Example



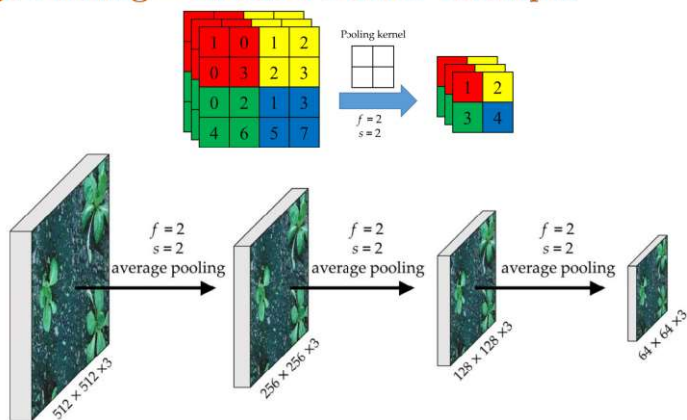
## POOLING LAYER

### • Average Pooling Results



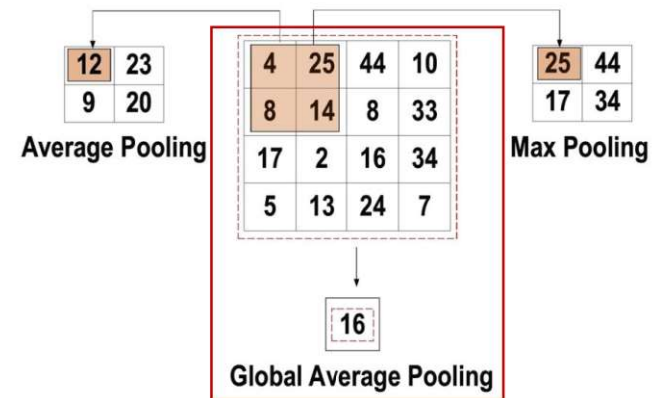
## POOLING LAYER

### • Average Pooling – Order 3 Tensor Example



## POOLING LAYER

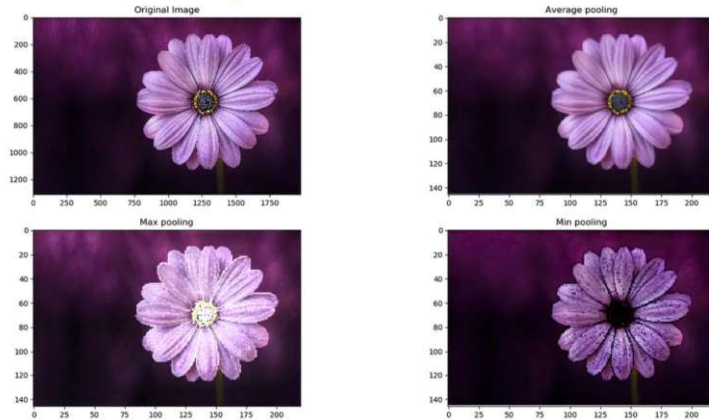
### • Global Average Pooling - Example





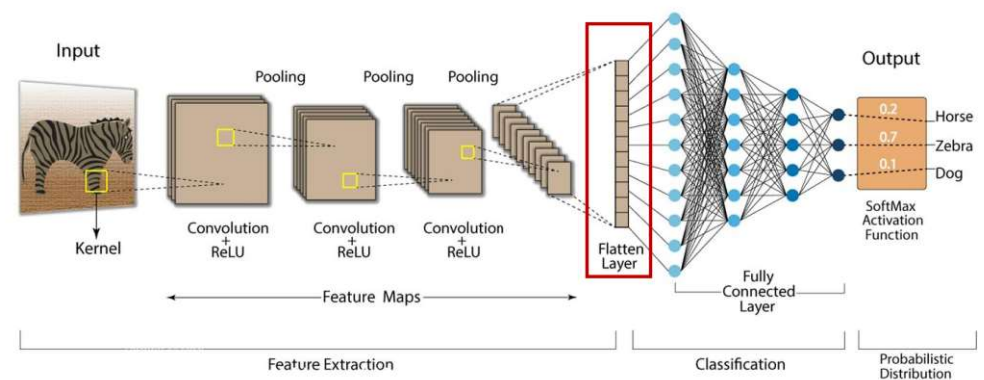
## POOLING LAYER

### • Some Further Examples



*Pooling layer not only effectively compresses the amount of data and parameters, reduces the feature map dimension, minimizes overfitting and reduces the sensitivity of the network to minor variations in input structure*

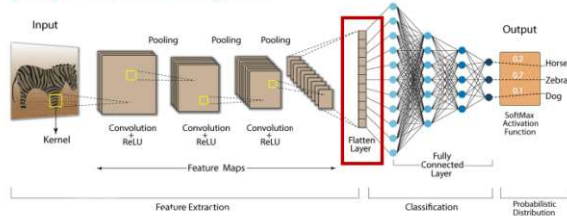
## CNN ARCHITECTURE FLATTENING LAYER



## FLATTENING LAYER

### • The Purpose

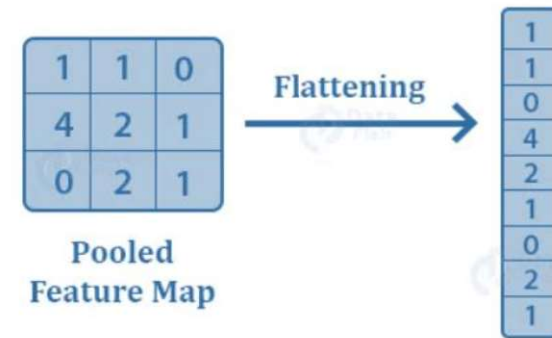
- The **FLATTEN LAYER** serves the purpose of *reshaping* the *output* of the preceding layer *into a one-dimensional vector*, which can then be *fed into* subsequent *fully connected layers*.



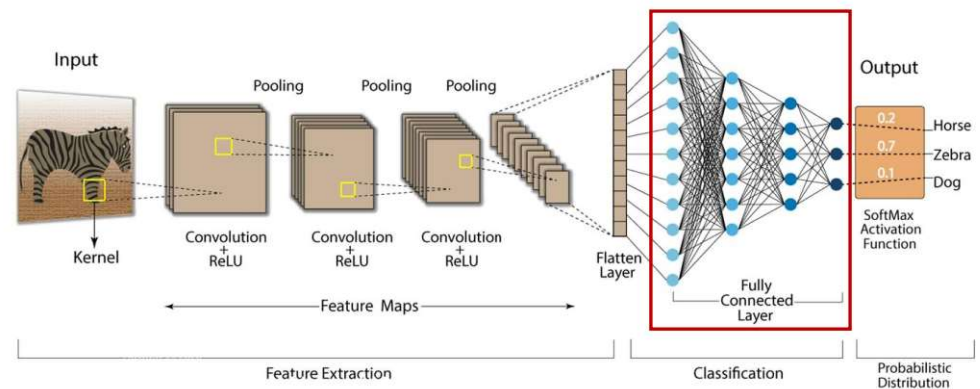
- This is *typically done* to prepare the data for input into a fully connected layer or another type of *layer that requires a one-dimensional input*.

## FLATTENING LAYER

### • How it Works?



## CNN ARCHITECTURE FULLY CONNECTED LAYER



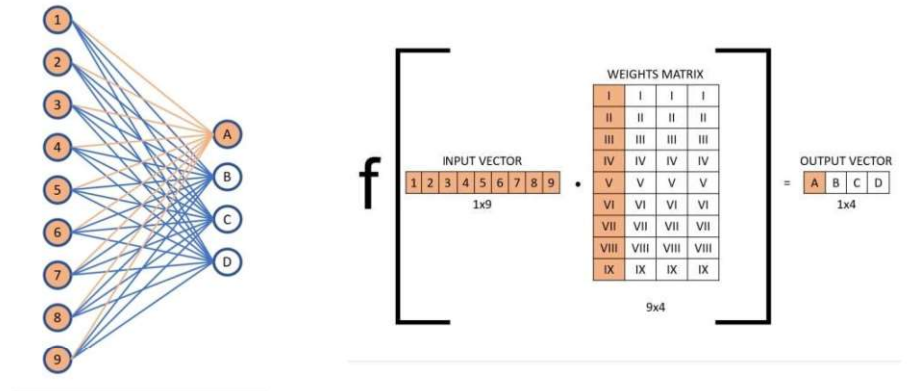
## FULLY CONNECTED LAYER

### • Introduction

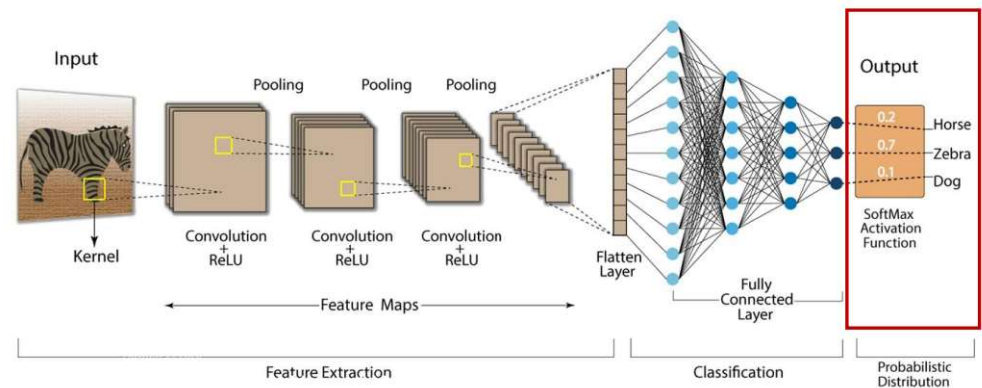
- The **FULLY CONNECTED** layers are *used to classify* the input data *based* on the *features extracted* by the convolutional and pooling layers.
- They are *typically the last layers* in the CNN architecture, and they are used to *produce the final output of the model*.
- The distilled *knowledge* of the last layer is *passed* through *fully-connected layers* with different hidden layers, followed by a *SoftMax activation* function.

## FULLY CONNECTED LAYER

### • How it Works?



## CNN ARCHITECTURE SOFTMAX ACTIVATION

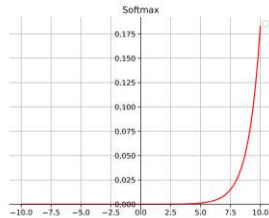


## SOFTMAX ACTIVATION

### • Introduction

- **SOFTMAX** is an *activation function* for *multi-classification problems*.

$$\text{Softmax}(x) = \frac{e^{x_i}}{\sum_{j=1}^K e^{x_j}}$$



- For any *real vector* of *length K*, *Softmax activation* can compress it into a *real vector* of *length K*, with values in the *range (0, 1)* and *vector elements summing to 1*.

## CODE EXAMPLE

## CODE EXAMPLE

### • Convolutional Neural Network (CNN)

- <https://www.tensorflow.org/tutorials/images/cnn>

## REFERENCES



# REFERENCES

1. <http://yann.lecun.com/exdb/mnist>
2. Introduction to Convolutional Neural Networks, Jianxin Wu
3. <https://cs231n.github.io/convolutional-networks/>
4. Convolutional neural networks: an overview and application in radiology, Rikiya Yamashita et al., 2018
5. <https://towardsdatascience.com/a-practical-guide-on-convolutional-neural-networks-cnns-with-keras-21421172005e>
6. Injurious or Noninjurious Defect Identification From MFL Images in Pipeline Inspection Using Convolutional Neural Network, Jian Feng et al., 2017
7. Review of deep learning: concepts, CNN architectures, challenges, applications, future directions, Laith Alzubaidi et al., 2021
8. An Introduction to Convolutional Neural Networks, Keiron O Shea et al., 2015
9. A review of convolutional neural networks in computer vision, Xia Zhao et al., 2024
10. <https://builtin.com/machine-learning/fully-connected-layer>
11. <https://pub.towardsai.net/introduction-to-pooling-layers-in-cnn-da661eab34>
12. RunPool: A Dynamic Pooling Layer for Convolution Neural Network, Putra Wanda et al., 2020
13. An Improved U-Net Model Based on Multi-Scale Input and Attention Mechanism: Application for Recognition of Chinese Cabbage and Weed, Zhongyang Ma et al., 2023